

# Simply Ansible

| Getting started with Ansible best practices

David M. Lee, II

[@leedm777](#) GitHub

[@leedm777](#) Twitter

[leedm777@yahoo.com](mailto:leedm777@yahoo.com)

## **new Ansible()**

Ansible is an IT automation tool. It can configure systems, deploy software, and orchestrate more advanced IT tasks such as continuous deployments or zero downtime rolling updates.

# Ansible basics

- Ansible uses **YAML** for pretty much everything
  - Except for the occasional **INI format**
- Templates are written in **Jinja2** template language
  - Can show up in `.yaml` files. Much fun!
- Requires Python 2.x, on both control machine and managed nodes
  - No agent required on managed nodes
  - Modules may require certain libraries/apps to be installed

# Using Ansible



## ansible - ad hoc commands

Handy, but use sparingly. Changes to systems should be done through playbooks.

```
$ ansible all \
> --module-name command \
> --args "uname -a"
services | SUCCESS | rc=0 >>
Linux services 3.2.0-107-virtual #148-Ubuntu SMP Mon Jul 18

mariadb | SUCCESS | rc=0 >>
Linux mariadb 3.2.0-107-virtual #148-Ubuntu SMP Mon Jul 18

pdns | SUCCESS | rc=0 >>
Linux pdns 3.10.0-327.22.2.el7.x86_64 #1 SMP Thu Jun 23 1
```

## ansible-playbook - what you really want


The `ansible-playbook` command executes the tasks specified in the playbook, in order, against the managed hosts.

```
$ ansible-playbook site.yml
PLAY [Do the thing] *****

TASK [setup] *****
ok: [mariadb]
ok: [services]
ok: [pdns]

TASK [The thing] *****
ok: [mariadb]
ok: [services]
skipped: [pdns]
```

# Anatomy of an Ansible project



# Inventory

## The list of managed hosts

Can be static (in a `.ini` file) or dynamic (results of a `.py` inventory script).

```
# ./inventory/vagrant.ini
[services]
services ansible_host=192.168.98.100

[mariadb]
mariadb ansible_host=192.168.98.200
```



# Variables

- Values that may vary per system
  - About a zillion different places to specify vars
  - Please don't use *all* of them

```
# ./group_vars/all/main.yml
---
# default to direct mail delivery.
# if you have a relay, set it here.
smtp_relay_host: ''

# Use for NODE_ENV, RAILS_ENV, RACK_ENV, etc.
app_env: development
```

# Secret Variables: Ansible Vault

The `ansible-vault` command can be used to manage encrypted files. Put any secrets for a given environment in a `vault.yml` file in that environment's `group_vars` .

```
# create a vault
$ ansible-vault create group_vars/${ENV}/vault.yml
# edit a vault
$ ansible-vault edit group_vars/${ENV}/vault.yml
```

# Playbook

A playbook is a list of plays, which are executed sequentially, in order.

```
# ./site.yml
---
- name: Some play
  hosts: some-hosts
  roles:
    - some-role
- name: Some other play
# ...
```

# Play

A play is a list of roles, which are executed sequentially, in order.

```
- name: Some play
  hosts: some-hosts
  roles:
    - some-role
    - some-other-role
```

# Role

A role is a list of tasks, which are executed sequentially, in order.  
Plus some other stuff.

```
site.yml
roles/
  some-role/
    tasks/
      main.yml # tasks go in here
    vars/
      main.yml # role-specific vars go in here
    handlers/
      main.yml # tasks that are triggered optionally
                # i.e. service restarts
    files/
      ... # content for copy/script tasks
    templates/
      ... # content for template tasks
          # always append .j2 to filenames
```

# Parameterized Role

A role may be parameterized, meaning that it expects certain variables to be set.

```
- role: rvm_io.rvm1-ruby  
  rvm1_autolib_mode: 4 # automatically install deps  
  rvm1_rubies: ['ruby-2.3.1', 'ruby-2.2.5']
```

# Sharing roles

- Common roles can be shared via [Ansible galaxy](#) or [git repo](#)
- BUT... they tend to be very small, poorly versioned, and so poorly maintained, this is rarely useful
  - Don't use a role unless it has tagged versions
  - AND it looks like someone loves and cares for it
  - Otherwise, learn what you can from it and write your own

# Task

A task specifies a module and the parameters to invoke it with. The task is executed in parallel on every matching host for that play.

- name: perform an action  
*# args can be name=value*  
action: some\_param=some-value
- name: perform another action  
*# or args can be an object*  
action:  
    some\_param: some-value



# Block

A block can allow you to more easily apply settings to a group of tasks. Also adds `rescue` blocks for error handling. New in Ansible 2.x.

```
- when: some_var == 'some-val'  
  block:  
    - name: conditionally do a thing  
      action: some_param=some-value  
    - name: conditionally do another thing  
      another_action: some_param=some-value
```

# Template

You can use [Jinja2 syntax](#) in `.j2` template files, or in task definitions.

```
- name: perform an action
  action: some_param='{{ some_var }}'
#               ^               ^
#               in YAML, quotes are
#               usually necessary to
#               avoid yaml/j2 confusion
```

# Tags and limits

**Tags** specify which tasks to run. **Limits** specifies which hosts to run those tasks on.

```
# run entire playbook on single machine
$ ansible-playbook site.yml --limit some-new-machine

# run everything tagged with sontaran on all machines
$ ansible-playbook site.yml --tags sontaran

# run sil and ood on the webservers not in San Diego
$ ansible-playbook site.yml --tags sil,ood \
>   --limit webservers:!san-diego
```

# PROTIPS



# Put tasks in roles, not the playbook

While you can list tasks directly in a play, you're usually better off putting the task off in a role.

```
# ./site.yml
- name: Some play
  hosts: some-hosts
  pre_tasks:          # meh
    - name: some pre_task
  roles:
    - some-role
  tasks:              # really?
    - name: some task
  post_tasks:         # PLEASE STOP!
    - name: some post_task
```

## Really, keep tasks in roles

Besides, content from Ansible Galaxy can only be brought in as a role, so trying to sort that into existing `{pre,post}_tasks` is a pain.

```
- name: Some play
  hosts: some-hosts
  pre_tasks:
    # but what if I wanted to call a galaxy role
    # before this?
    - name: some pre_task
  post_tasks:
    - name: some post_task
    # or maybe a role after this?
```

## Keep roles small and single purpose

If the playbook is your 'program', then roles are your 'functions'.

- For example, `pip` , `firewall` , `nodejs` , `timezone` , `auton`
- Not things like `common` , `general` , `misc` , ...
- Keep them small; most will fit on the screen

## Bending the rule: debug-tools

I usually have a rule with a bunch of one-line installers for things that are generally useful for debugging. But if you can't install it with a single `apt` , `pip` , `npm` , `gem` , etc., then it gets its own role.

```
# ./roles/debug-tools/tasks/main.yml
- when: ansible_distribution == 'Ubuntu'
  name: apt-get install <cool-stuff>
  apt: name={{ item }}
  with_items:
    - curl
    - htop
    - iotop
    - jq
  # ...
```



## Avoid repeating roles...

Here both `foo` and `bar` need `zygon`. Since plays are sequential, `zygon` will be installed in sequence. Can really slow down the playbook.

```
- hosts: foo
  roles:
    - zygon # Bad!
    - foo

- hosts: bar
  roles:
    - zygon # Bad!
    - bar
```

## i.e., consolidate roles to get parallelism

Since the `zygon` role is in a single play, it will execute in parallel on both `foo` and `bar`. You can't always do this, but when you can, it helps tremendously.

```
- hosts: foo,bar # both foo and bar groups
  roles:
    - zygon      # installs in parallel
                  # across foo & bar nodes

- hosts: foo
  roles:
    - foo

- hosts: bar
  roles:
    - bar
```

## Tag each role with the role's name

By keeping tag names consistent with role names, it's easier to keep track of what tags you have.

```
# ./roles/zygon/tasks/main.yml
---
- tags: zygon
  block: # Ansible 2.x syntax
    - name: First zygon invasion
      # ...
```

Since roles are small and single purpose, you're usually debugging individual roles.

```
$ ansible-playbook site.yml --tags zygon
```

## Role tagging in Ansible 1.9

Before Ansible 2.x, the best way to tag an entire role was to put the tag in the playbook. Either way is fine.

```
roles:  
- { role: zygon, tags: zygon } # Ansible 1.9 syntax
```

# Don't put tons of tags on a task

If you find yourself putting lots of tags on a task, you're probably doing it wrong. A tag should be there for a reason.

```
tags:
```

- provision *# okay, but every task in the  
# playbook is for provisioning*
- bamboo *# and everything is for bamboo*
- bamboo-agent  
*# what are you doing?*
- bamboo-agent-ubuntu  
*# wait...*
- bamboo-agent-ubuntu-precise  
*# MAKE IT STOP*
- bamboo-agent-ubuntu-precise-provision  
*# \*sigh\**

## Put variable definitions in the right spot

```
site.yml          # in the occasional vars block
                  # for parameterized roles

group_vars/
  all/
    main.yml      # default values
    ...           # can break into multiple files
  staging/        # Have a group for each environment
    main.yml      # override defaults for env
    vault.yml     # safe place for secrets
roles/
  kandy-man/
    vars/
      main.yml    # specific vars for the role
                  # version numbers, SHA hashes, etc.
```

# Vaults: too many vaults will kill you

- We tried having single-var per vault
  - PRO: Avoids merge conflicts w/ vaults
  - CON: No one follows the rules
    - Vars end up in random vaults
    - Or var was renamed and vault wasn't
    - Grep is either powerless, or really slow
- Use [ansible-vault-tools](#) and as few vaults as possible
  - Single vault per `group_vars`
  - `ansible-vault-merge` - magical vault merging
  - `ansible-vault-textconv` - magical vault diffing
    - And `git grep --textconv` can search in vaults!

# !!!ENCRYPT YOUR VAULT PASSWORD!!!

The `ansible-vault-tools` project also provides `gpg-vault-password-file`, which can use your PGP key to encrypt your vault password.

- You do [have a PGP key](#), right?
- And it's [encrypted with a passphrase](#)?
- And you use [gpg-agent](#) so you don't go insane typing the passphrase all the time?



# Put common settings in `ansible.cfg`


Defaults should be what's best for the humans that run the playbooks by hand the most (probably devs).

```
# ./ansible.cfg
[defaults]
# Why is the default not smart?
gathering = smart
# Default for most common params devs would use
inventory = ./inventory/vagrant.ini
# Consistent configuration for vault passwords
vault_password_file = $HOME/.ansible/some_project_vault

[privilege_escalation]
# Nearly everything requires sudo, so default to it
become = True

[ssh_connection]
# Speeds things up, but requires disabling
# requiretty in /etc/sudoers
pipelining = True
```

# Make sure roles are idempotent



The concept that change commands should only be applied when they need to be applied, and that it is better to describe the desired state of a system than the process of how to get to that state.

-- [Ansible Glossary](#)

# Handlers are not idempotent!

...but often necessary. Just be aware of this when a playbook fails.

```
# ./site.yml
- # ...
  roles: [ master, cyberman ]

# ./roles/master/tasks/main.yml
- cp: src=missy.conf dest=/etc/master/
  notify: restart master
  # ^^ schedules restart of master if file is copied

# ./roles/cyberman/tasks/main.yml
- fail: msg="***HEAD EXPLODES***"
  # ^^ causes play to fail before restart master
  # even if you fix and re-run, file already copied
  # so it won't notify on re-run
```

# Use a wrapper script so everyone runs the playbook consistently

- Setup `virtualenv` so everyone runs the same version
- Add some common/useful options ( `-v` , `--diff` )
- Add domain specific options, useful for your team
  - `--env` to specify environment
  - `--service` to deploy a specific service
  - `--` to break for `ansible-playbook` options

```
$ ./deploy.sh --env staging -- --tags cyberman --check
+ exec ansible-playbook -v --diff \
+   --inventory-file ./inventory/staging.ini \
+   --tags cyberman \
+   --check \
+   -- site.yml
Using /Users/dlee/prj/doctor/ansible.cfg as config file

PLAY [pre-roles] *****
```

# Don't make a role portable until you have to

There is no such thing as portable code. Only code that's been ported.

Avoid supporting multiple platforms as much as you can. It's better to standardize on a distro/version.

That said,

- You usually end up with one weird thing that needs to be on another distro
- Or you'll have a transition time as you upgrade major versions
- Don't port everything; only port what you need; one thing at a time.

# Logic goes in roles; not in playbooks

Ansible allows you to split logic between playbooks and roles, but that just makes it harder to figure out what's going on.

```
# ./site.yml
- # ...
  roles:
    - role: epel
      when: ansible_distribution == 'CentOS' # bad

# ./roles/epel/tasks/main.yml
- when: ansible_distribution == 'CentOS' # Good!
  block:
    - name: yum install epel-...
```

# Test your playbooks

- A `make test` target is always useful
- `ansible-playbook --syntax-check` finds syntax errors
- `ansible-lint` finds common errors

```
test: virtualenv galaxy
    @$(MAKE) check-env # errors if we're
                        # not in virtualenv
    ansible-playbook site.yml --syntax-check
    ansible-lint --exclude build/ site.yml
.PHONY: test
```

# Questions





**Thanks!**

